

Language Models in Cryptanalysis

How do we crack the long ciphers?

Morten Munk

November 2025

AAU CPH - SW9

Contents

1.	The papers	2
1.1	Why these papers?	4
2.	Background	5
2.1	Recap from last time	6
2.2	Standard Attention Computation	8
3.	Linformer	10
4.	Performer	14
5.	Reordered Computation	17
6.	Long Short Transformer	21
7.	Comparative study	24
8.	Limitations	27
9.	Summary	29
10.	Questions?	31

1. The papers

1. The papers

X-former Elucidator: Reviving Efficient Attention for Long Context Language Modeling

Miao, et al., 2024

Transformers are RNNs: Fast Autoregressive Transformers with Linear Attention

Katharopoulos, et al., 2020

Long-Short Transformer: Efficient Transformers for Language and Vision

Zhu, et al., 2021

Rethinking Attention with Performers

Choromanski, et al., 2021

Linformer: Self-Attention with Linear Complexity

Wang, et al., 2020

1.1 Why these papers?

Efficient computation during training

- Causal LM inference is fast

Best methods from comparative study

- We don't have time for them all
- Some are more inference focused

Summary

- We care about efficiency during training
- Inference for long cipher struggle due to lack of generalization on long ciphers

2. Background

2.1 Recap from last time

Homophonic Substitution Ciphers

- 1:>0 mappings
- English without spaces & punctuation




Figure 1: Example of a homophonic substitution cipher. (Kambhatla et al., Findings 2023)

2.1 Recap from last time

Causal LM

- Learns both cipher & plaintext
- Reads left to right

Seq2seq

- Learns plaintext only
- Bidirectional

Both suffer from $O(N^2)$ attention 😞

#keys	Model	Max Len.	
		400	700
30-45	Seq-to-Seq	72.30	fail
	PrefixLM	54.73	69.50
	CausalLM (tgt)	29.99	37.20
	CausalLM	0.40	0.21
40-65	PrefixLM	69.50	54.73
	CausalLM (tgt)	29.99	37.20
	CausalLM	0.83	0.80
30-85	PrefixLM	70.52	71.82
	CausalLM (tgt)	42.05	42.69
	CausalLM	2.25	2.19

Figure 2: SER on synthetic HS ciphers.
(Kambhatla et al., Findings 2023)

2.2 Standard Attention Computation

Rows (Queries)

- Token we are looking for

Columns (Keys)

- Token we are looking at

Values (Cells)

- Attention Score

Notice: $(N \times N) = N^2$

Cipher: X Y Z

$$X \rightarrow [Y, Z]$$

$$Y \rightarrow [X, Z]$$

$$Z \rightarrow [X, Y]$$

	X	Y	Z
X	$X \rightarrow X$	$X \rightarrow Y$	$X \rightarrow Z$
Y	$Y \rightarrow X$	$Y \rightarrow Y$	$Y \rightarrow Z$
Z	$Z \rightarrow X$	$Z \rightarrow Y$	$Z \rightarrow Z$

Table 1: Attention Weight Matrix ($N \times N$)

2.2 Standard Attention Computation

Why is $O(N^2)$ Ineffecient?

Many attention heads

- Each head has its own matrix

Many forward passes

- Each matrix recomputed at each pass

Causal LM as an example

- 12 layers \times 12 heads =
144 attention heads

**Ciphers with 1000s of characters
does not scale well 😭**




Figure 3: Example of attention with multiple heads. (Yuan et al., 2022)

3. Linformer

3. Linformer

Disclaimer: Math is heavily simplified for understanding !

Standard attention

$$\text{attention}(Q, K, V) = \underbrace{\text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right)}_P V$$

Claim: Attention matrix is low-rank

- It can be represented by a smaller matrix

For any Q, K, V exists a low-rank matrix \tilde{P} where:

- \tilde{P} has minimal error
- \tilde{P} has low-rank (fewer dims/features)

3. Linformer

Disclaimer: Math is heavily simplified for understanding !

Standard attention

$$\text{attention}(Q, K, V) = \underbrace{\text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right)}_P V$$

Claim: Attention matrix is low-rank

- It can be represented by a smaller matrix

For any Q, K, V exists a low-rank matrix \tilde{P} where:

- \tilde{P} has minimal error
- \tilde{P} has low-rank (fewer dims/features)

I will skip mathematical proof, and show empirical proof !

3. Linformer

Eigenvalue index

- Top 128 eigenvalues are most important
- Trailing 384 eigenvalues are not so important

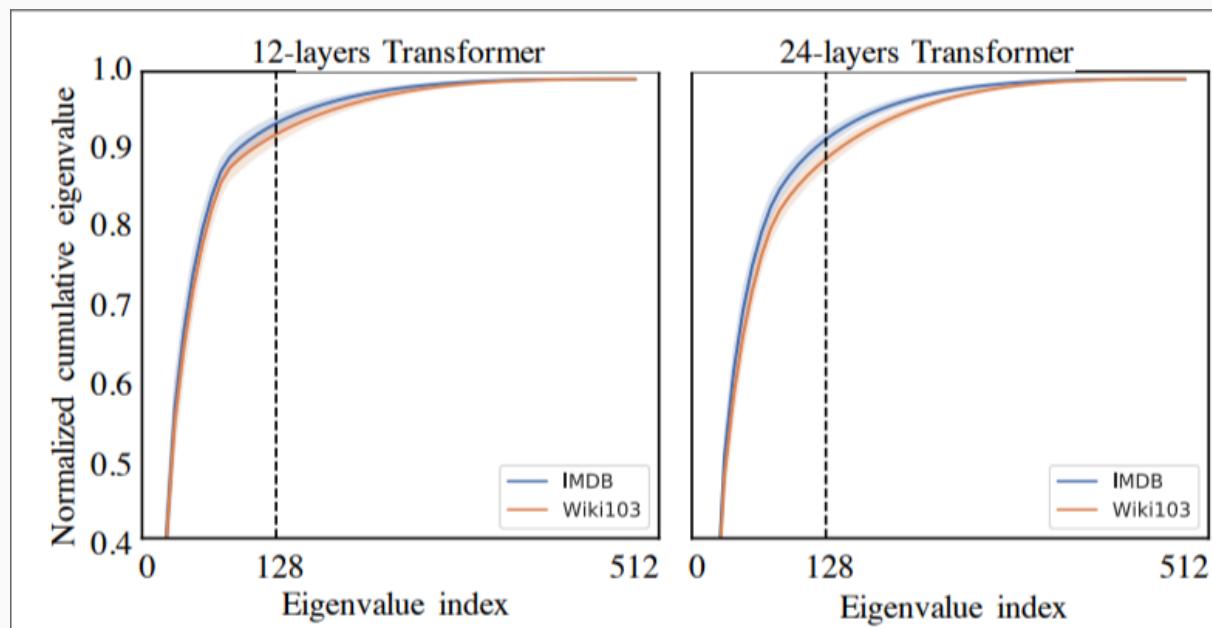


Figure 4: RoBERTa, IMDB & Wiki103 (Wang, et al., 2020)

3. Linformer

Similar to standard transformer

- Notice projections!

$$\text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right)V \rightarrow \text{softmax}\left(\frac{Q(E_K K)^T}{\sqrt{d}}\right)E_V V$$

$$n \times n = O(n^2)$$

- inefficient!

$$n \times k = O(nk)$$

- Better because $k \ll n$




Figure 5: (Wang, et al., 2020)

4. Performer

4. Performer

FAVOR+: Fast Attention Via positive Orthogonal Random Features

Linformer compresses - Performer approximates

Left side (Standard):

- Each key looks at every other key $O(L^2)$
- Exponential similarity e^{QK^T}
 - Grows positively/negatively based on similarity

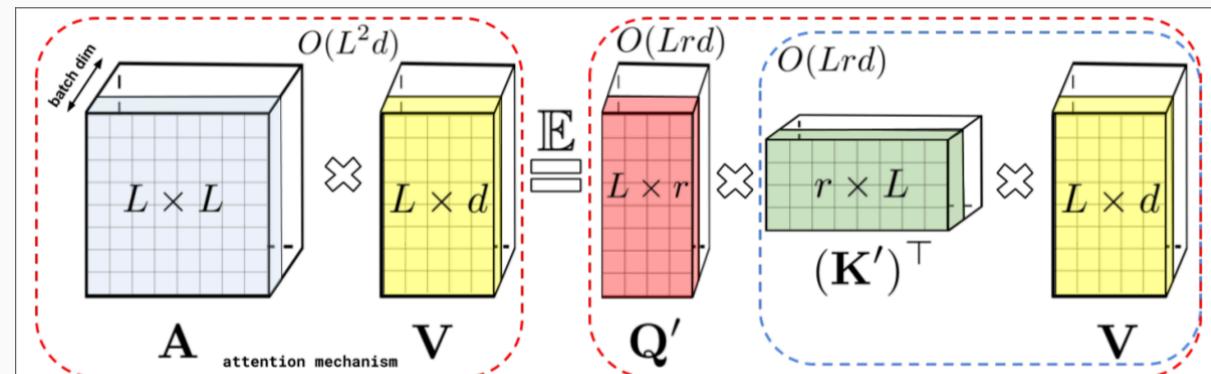


Figure 6: (Choromanski, et al., 2021)

4. Performer

Right side (Performer):

- Kernel trick with random features $\varphi(Q) \cdot \varphi(K) \approx e^{q_i \cdot k_j}$
- Same exponential similarity but smaller matrices

$$\varphi(Q)(\varphi(K)^T V) \rightarrow (L \times r) \times (r \times d) = O(Lr)$$

- r parameter determine number of directions to mimic similarities
 - Large r = more accurate, but slower




Figure 7: (Choromanski, et al., 2021)

5. Reordered Computation

5. Reordered Computation

Transformers are RNNs...

5. Reordered Computation

Transformers are RNNs...

- Not really 😊

5. Reordered Computation

Transformers are RNNs...

- Not really 

But this paper shows how they can act like it

Also it has no illustrations, so instead you get math 

5. Reordered Computation

Generalized Attention $O(N^2)$

- Output is weighted matrix determined by similarity function
- Sim() could be softmax or cosine similarity, etc.

$$V'_i = \frac{\sum_{j=1}^N \text{sim}(Q_i, K_j) V_j}{\sum_{j=1}^N \text{sim}(Q_i, K_j)}$$

Linear Attention $O(N)$

- Notice the reordered computation
- Uses rule of associativity $(A \times B) \times C = A \times (B \times C)$

$$V'_i = \frac{\varphi(Q_i)^T \sum_{j=1}^N \varphi(K_j) V_j^T}{\varphi(Q_i)^T \sum_{j=1}^N \varphi(K_j)}$$

5. Reordered Computation

Generalized Attention $O(N^2)$

- Output is weighted matrix determined by similarity function
- Sim() could be softmax or cosine similarity, etc.

$$V'_i = \frac{\sum_{j=1}^N \text{sim}(Q_i, K_j) V_j}{\sum_{j=1}^N \text{sim}(Q_i, K_j)}$$

Linear Attention $O(N)$

- Notice the reordered computation
- Uses rule of associativity $(A \times B) \times C = A \times (B \times C)$

$$V'_i = \frac{\varphi(Q_i)^T \sum_{j=1}^N \varphi(K_j) V_j^T}{\varphi(Q_i)^T \sum_{j=1}^N \varphi(K_j)} \rightarrow (\underbrace{\varphi(Q)\varphi(K)^T}_{M_1} V) = \underbrace{(\varphi(Q)\underbrace{(\varphi(K)^T V)}_{M_2})}_{O(N)}$$

Calculate M_2 which is independent of Att matrix size, then multiply with M_1

Softmax is non-factorizable! - replace with positive factorizable ELU allows reordering ✓

5. Reordered Computation

But why do they claim transformers are RNNs?

5. Reordered Computation

But why do they claim transformers are RNNs?

Used during inference

Standard Attention

- For each Q , look at each K

RNN approach

- s_i = content memory of keys and values
- z_i = normalizer memory (running total of weights)
- Combined they keep weighted average for next output
 - ▶ no need to look back at individual past keys!

For each Q it is constant lookup!

6. Long Short Transformer

6. Long Short Transformer

Linear complexity from dual attention

Short term (local window)

- Fine-grained local correlations
- Divide input sequence into w sized segments
- Sliding window looks at home segment and $\frac{w}{2}$ tokens on both sides
- Fixed size keeps attention $O(N)$

Long range (dynamic projection)

- Distant correlations across entire sequence
- Dynamic Projection P decides which keys/values are important to keep
- P creates low-rank version of K and V
- Low-rank contains fixed r summary points

6. Long Short Transformer

Aggregation of attention Final output is calculated at every head

- Query looks at concatenation of global and local keys/values

Scale mismatch Bias towards short term due to larger values

- Solved with DualLN normalization technique

Runtime depends on sequence length N and summary points r

- Hence linear runtime $O(N)$

7. Comparative study

7. Comparative study

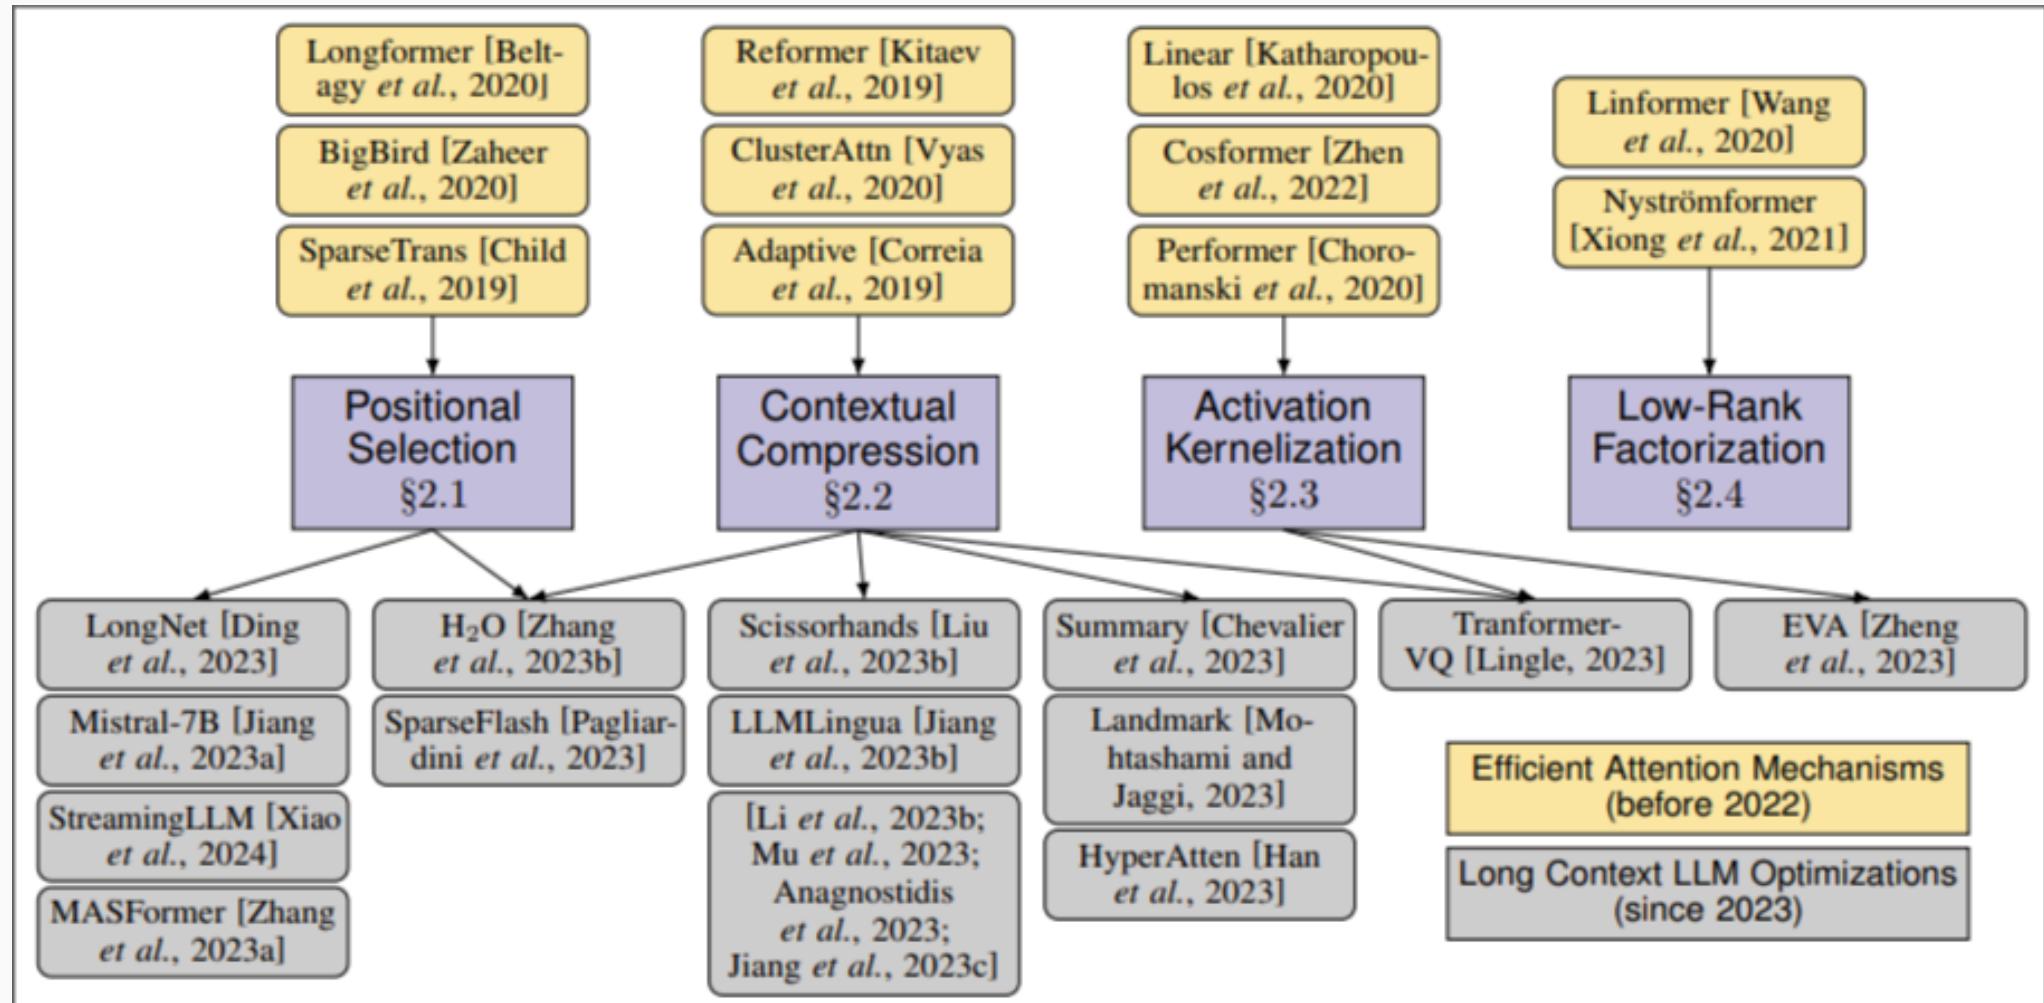


Figure 8: (Miao, et al., 2024)

7. Comparative study




Figure 9: (Miao, et al., 2024)

8. Limitations

8. Limitations

None of the methods are tested on ciphers

- Ciphers are slightly different than regular machine translation

No one method dominates

- Performer is likely best for 4096 sequence length
- I would like to go way beyond 4k sequence length 😎

Long Short dependencies

- Will it work for very long ciphers?
- What do you think is most important?
 - Global
 - Local
 - Evenly?

9. Summary

9. Summary

Linformer

- Low-rank projection

Performer

- Random feature approximation

Reordered

- Associativity method

Long Short

- Dual attention

10. Questions?

Please no math questions 😞