

Language Models in Cryptanalysis

How do we crack the long ciphers?

Morten Munk

November 2025

AAU CPH - SW9

Contents

1. The papers	2
1.1 Why these papers?	4
2. Background	5
2.1 Recap from last time	6
2.2 Standard Attention Computation	8
3. Linformer	10
4. Performer	14
5. Reordered Computation	17
6. Questions?	21

1. The papers

1. The papers

X-former Elucidator: Reviving Efficient Attention for Long Context Language Modeling

Miao, et al., 2024

Transformers are RNNs: Fast Autoregressive Transformers with Linear Attention

Katharopoulos, et al., 2020

Long-Short Transformer: Efficient Transformers for Language and Vision

Zhu, et al., 2021

Rethinking Attention with Performers

Choromanski, et al., 2021

Linformer: Self-Attention with Linear Complexity

Wang, et al., 2020

1.1 Why these papers?

Efficient computation during training

- Causal LM inference is fast

Best methods from comparative study

- We don't have time for them all
- Some are more inference focused

Summary

- We care about efficiency during training
- Inference for long cipher struggle due to lack of generalization on long ciphers

2. Background

2.1 Recap from last time

Homophonic Substitution Ciphers

- 1:>0 mappings
- English without spaces & punctuation

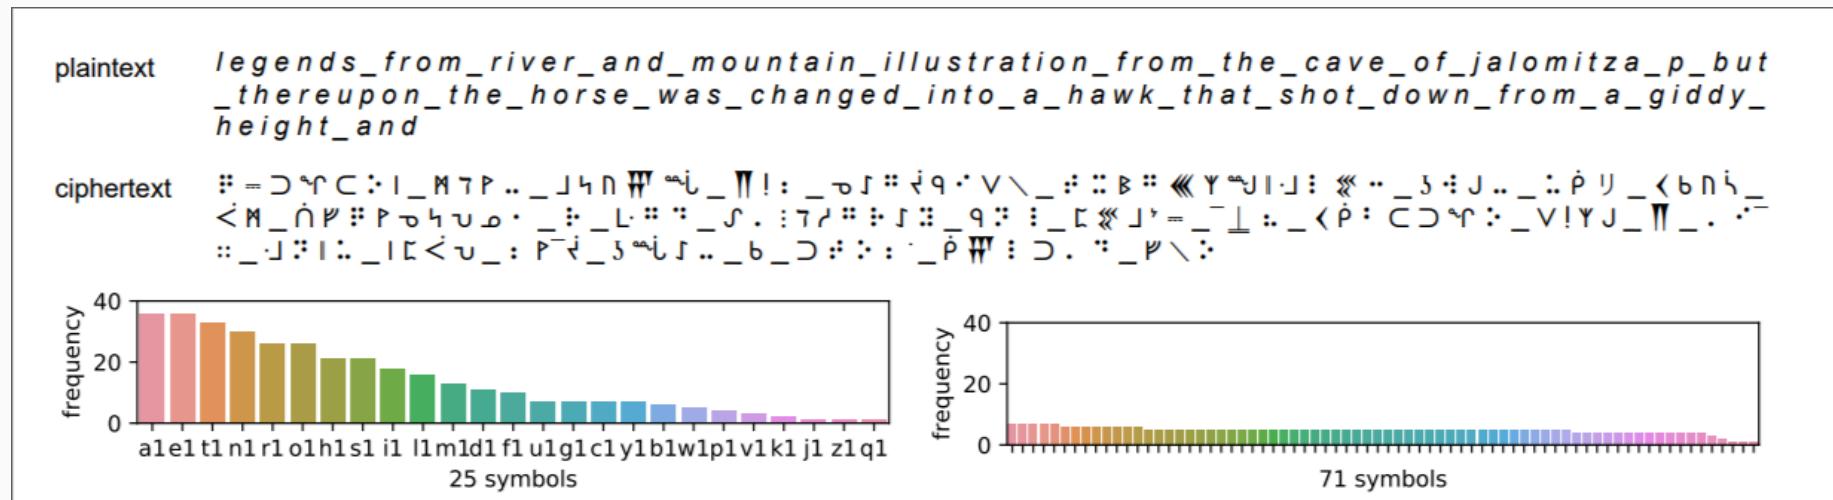


Figure 1: Example of a homophonic substitution cipher. (Kambhatla et al., Findings 2023)

2.1 Recap from last time

Causal LM

- Learns both cipher & plaintext
- Reads left to right

Seq2seq

- Learns plaintext only
- Bidirectional

Both suffer from $O(N^2)$ attention 😞

#keys	Model	Max Len.	
		400	700
30-45	Seq-to-Seq	72.30	fail
	PrefixLM	54.73	69.50
	CausalLM (tgt)	29.99	37.20
	CausalLM	0.40	0.21
40-65	PrefixLM	69.50	54.73
	CausalLM (tgt)	29.99	37.20
	CausalLM	0.83	0.80
30-85	PrefixLM	70.52	71.82
	CausalLM (tgt)	42.05	42.69
	CausalLM	2.25	2.19

Figure 2: SER on synthetic HS ciphers.
(Kambhatla et al., Findings 2023)

2.2 Standard Attention Computation

Rows (Queries)

- Token we are looking for

Columns (Keys)

- Token we are looking at

Values (Cells)

- Attention Score

Notice: $(N \times N) = N^2$

Cipher: X Y Z

$$X \rightarrow [Y, Z]$$

$$Y \rightarrow [X, Z]$$

$$Z \rightarrow [X, Y]$$

	X	Y	Z
X	$X \rightarrow X$	$X \rightarrow Y$	$X \rightarrow Z$
Y	$Y \rightarrow X$	$Y \rightarrow Y$	$Y \rightarrow Z$
Z	$Z \rightarrow X$	$Z \rightarrow Y$	$Z \rightarrow Z$

Table 1: Attention Weight Matrix ($N \times N$)

2.2 Standard Attention Computation

Why is $O(N^2)$ Ineffecient?

Many attention heads

- Each head has its own matrix

Many forward passes

- Each matrix recomputed at each pass

Causal LM as an example

- 12 layers \times 12 heads =
144 attention heads

**Ciphers with 1000s of characters
does not scale well 😭**

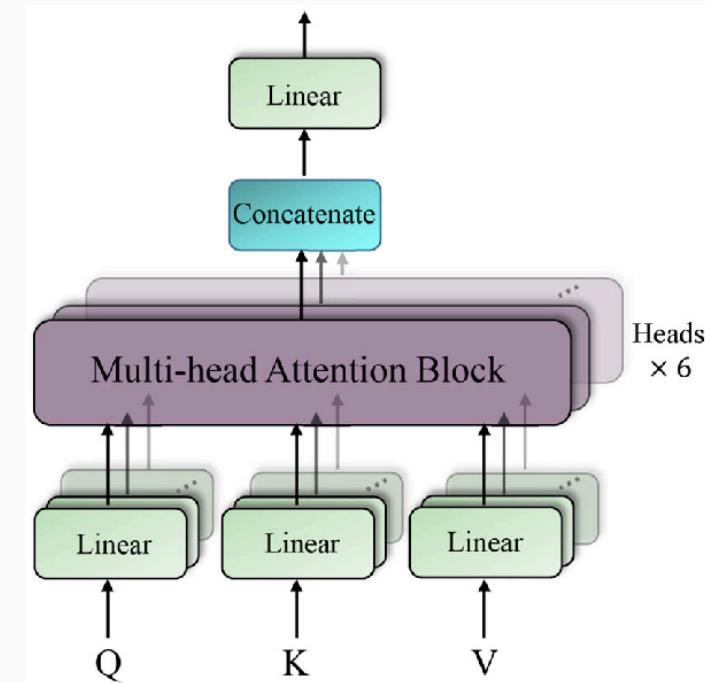


Figure 3: Example of attention with multiple heads. (Yuan et al., 2022)

3. Linformer

3. Linformer

Disclaimer: Math is heavily simplified for understanding !

Standard attention

$$\text{attention}(Q, K, V) = \underbrace{\text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right)}_P V$$

Claim: Attention matrix is low-rank

- It can be represented by a smaller matrix

For any Q, K, V exists a low-rank matrix \tilde{P} where:

- \tilde{P} has minimal error
- \tilde{P} has low-rank (fewer dims/features)

3. Linformer

Disclaimer: Math is heavily simplified for understanding !

Standard attention

$$\text{attention}(Q, K, V) = \underbrace{\text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right)}_P V$$

Claim: Attention matrix is low-rank

- It can be represented by a smaller matrix

For any Q, K, V exists a low-rank matrix \tilde{P} where:

- \tilde{P} has minimal error
- \tilde{P} has low-rank (fewer dims/features)

I will skip mathematical proof, and show empirical proof !

3. Linformer

Eigenvalue index

- Top 128 eigenvalues are most important
- Trailing 384 eigenvalues are not so important




Figure 4: RoBERTa, IMDB & Wiki103 (Wang, et al., 2020)

3. Linformer

Similar to standard transformer

- Notice projections!

$$\text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right)V \rightarrow \text{softmax}\left(\frac{Q(E_K K)^T}{\sqrt{d}}\right)E_V V$$

$$n \times n = O(n^2)$$

- inefficient!

$$n \times k = O(nk)$$

- Better because $k \ll n$




Figure 5: (Wang, et al., 2020)

4. Performer

4. Performer

FAVOR+: Fast Attention Via positive Orthogonal Random Features

Linformer compresses - Performer approximates

Left side (Standard):

- Each key looks at every other key $O(L^2)$
- Exponential similarity e^{QK^T}
 - Grows positively/negatively based on similarity

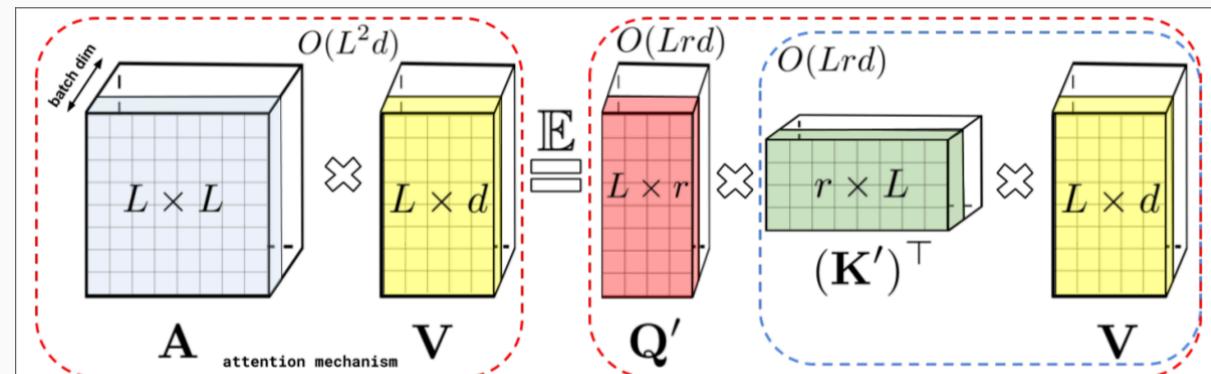


Figure 6: (Choromanski, et al., 2021)

4. Performer

Right side (Performer):

- Kernel trick with random features $\varphi(Q) \cdot \varphi(K) \approx e^{q_i \cdot k_j}$
- Same exponential similarity but smaller matrices

$$\varphi(Q)(\varphi(K)^T V) \rightarrow (L \times r) \times (r \times d) = O(Lr)$$

- r parameter determine number of directions to mimic similarities
 - Large r = more accurate, but slower




Figure 7: (Choromanski, et al., 2021)

5. Reordered Computation

5. Reordered Computation

Transformers are RNNs...

5. Reordered Computation

Transformers are RNNs...

- Not really 😊

5. Reordered Computation

Transformers are RNNs...

- Not really 

But this paper shows how they can act like it

Also it has no illustrations, so instead you get math 

5. Reordered Computation

Generalized Attention $O(N^2)$

- Calculates output of a query by weighted average of all V vectors
- Similarity between Q and K determines weight

$$V'_i = \frac{\sum_{j=1}^N \text{sim}(Q_i, K_j) V_j}{\sum_{j=1}^N \text{sim}(Q_i, K_j)}$$

Linear Attention $O(N)$

- Notice the reordered computation
- Uses rule of associativity $(A \times B) \times C = A \times (B \times C)$

$$V'_i = \frac{\varphi(Q_i)^T \sum_{j=1}^N \varphi(K_j) V_j^T}{\varphi(Q_i)^T \sum_{j=1}^N \varphi(K_j)}$$

5. Reordered Computation

Generalized Attention $O(N^2)$

- Calculates output of a query by weighted average of all V vectors
- Similarity between Q and K determines weight

$$V'_i = \frac{\sum_{j=1}^N \text{sim}(Q_i, K_j) V_j}{\sum_{j=1}^N \text{sim}(Q_i, K_j)}$$

Linear Attention $O(N)$

- Notice the reordered computation
- Uses rule of associativity $(A \times B) \times C = A \times (B \times C)$

$$V'_i = \frac{\varphi(Q_i)^T \sum_{j=1}^N \varphi(K_j) V_j^T}{\varphi(Q_i)^T \sum_{j=1}^N \varphi(K_j)} \rightarrow (\underbrace{\varphi(Q)\varphi(K)^T}_{M_1} \underbrace{V}_{M_2}) = \underbrace{\varphi(Q)}_{M_1} \underbrace{(\varphi(K)^T V)}_{M_2 O(N)}$$

Calculate M_2 which is independent of Att matrix size, then multiply with M_1

$N(O(N)) \rightarrow \text{linear!}$

5. Reordered Computation

But why do they claim transformers are RNNs?

5. Reordered Computation

But why do they claim transformers are RNNs?

Used during inference

Standard Attention

- For each Q , look at each K

RNN approach

- s_i = content memory of keys and values
- z_i = normalizer memory (running total of weights)
- Combined they keep weighted average for next output
 - ▶ no need to look back at individual past keys!

For each Q it is constant lookup!

6. Questions?
