Language Models in Cryptanalysis

How do we crack the long ciphers?

Morten Munk

November 2025

AAU CPH - SW9

Contents

1. The papers

X-former Elucidator: Reviving Efficient Attention for Long Context Language Modeling
Miao, et al., 2024

Transformers are RNNs: Fast Autoregressive Transformers with Linear Attention
Katharopoulos, et al., 2020

Long-Short Transformer: Efficient Transformers for Language and Vision
Zhu, et al., 2021

Rethinking Attention with Performers
Choromanski, et al., 2021

Linformer: Self-Attention with Linear Complexity
Wang, et al., 2020

Efficient computation during training

- Causal LM inference is fast

Best methods from comparative study

- We don’t have time for them all
- Some are more inference focused

Summary

- We care about efficiency during training
- Inference for long cipher struggle due to lack of generalization on long ciphers

2. Background

Homophonic Substitution Ciphers

- 1:>0 mappings
- English without spaces & punctuation

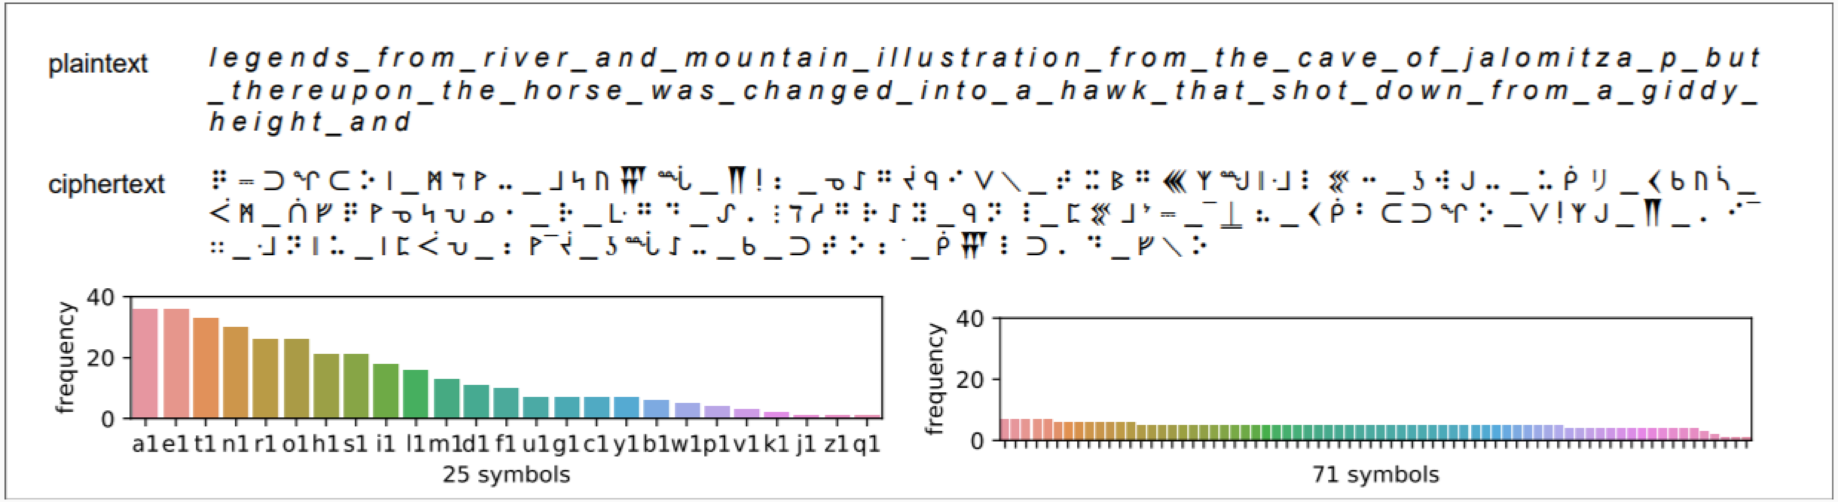

Figure 1: Example of a homophonic substitution cipher. (Kambhatla et al., Findings 2023)

Causal LM

- Learns both cipher & plaintext
- Reads left to right

Seq2seq

- Learns plaintext only
- Bidirectional

Both suffer from $O(N^2)$ attention 😔

#keys	Model	Max Len. 400	700
30-45	Seq-to-Seq	72.30	fail
	PrefixLM	54.73	69.50
	CausalLM (tgt)	29.99	37.20
	CausalLM	**0.40**	**0.21**
40-65	PrefixLM	69.50	54.73
	CausalLM (tgt)	29.99	37.20
	CausalLM	**0.83**	**0.80**
30-85	PrefixLM	70.52	71.82
	CausalLM (tgt)	42.05	42.69
	CausalLM	**2.25**	**2.19**

Figure 2: SER on synthetic HS ciphers. (Kambhatla et al., Findings 2023)

Rows (Queries)

- Token we are looking for

Columns (Keys)

- Token we are looking at

Values (Cells)

- Attention Score

Notice: $(N \times N) = N^2$

Cipher: X Y Z

$$X \rightarrow [Y, Z]$$

$$Y \rightarrow [X, Z]$$

$$Z \rightarrow [X, Y]$$

	X	**Y**	**Z**
X	$X \rightarrow X$	$X \rightarrow Y$	$X \rightarrow Z$
Y	$Y \rightarrow X$	$Y \rightarrow Y$	$Y \rightarrow Z$
Z	$Z \rightarrow X$	$Z \rightarrow Y$	$Z \rightarrow Z$

Table 1: Attention Weight Matrix $(N \times N)$

Why is $O(N^2)$ Ineffecient?

Many attention heads

- Each head has its own matrix

Many forward passes

- Each matrix recomputed at each pass

Causal LM as an example

- 12 layers $\times$ 12 heads $=$ 144 attention heads

Ciphers with 1000s of characters does not scale well 😭

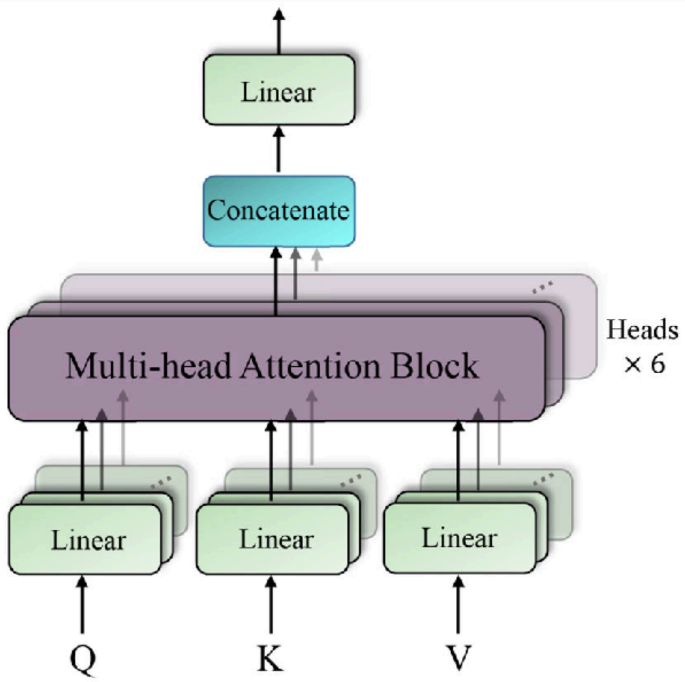

Figure 3: Example of attention with multiple heads. (Yuan et al., 2022)

3. Linformer

Disclaimer: Math is heavily simplified for understanding ⚠️

Standard attention

$$\text{attention}(\mathrm{Q}, \mathrm{K}, \mathrm{V}) = \underbrace{\text{softmax}\left(\frac{\mathrm{QK}^T}{\sqrt{d}}\right)}_{P} V$$

Claim: Attention matrix is low-rank

- It can be represented by a smaller matrix

For any Q, K, V exists a low-rank matrix $\tilde{P}$ where:

- $\tilde{P}$ has minimal error
- $\tilde{P}$ has low-rank (fewer dims/features)

Disclaimer: Math is heavily simplified for understanding ⚠️

Standard attention

$$\text{attention}(\mathrm{Q}, \mathrm{K}, \mathrm{V}) = \underbrace{\text{softmax}\left(\frac{\mathrm{QK}^T}{\sqrt{d}}\right)}_{P} V$$

Claim: Attention matrix is low-rank

- It can be represented by a smaller matrix

For any Q, K, V exists a low-rank matrix $\tilde{P}$ where:

- $\tilde{P}$ has minimal error
- $\tilde{P}$ has low-rank (fewer dims/features)

I will skip mathematical proof, and show empirical proof ⚠️

Eigenvalue index

- Top 128 eigenvalues are most important
- Trailing 384 eigenvalues are not so important

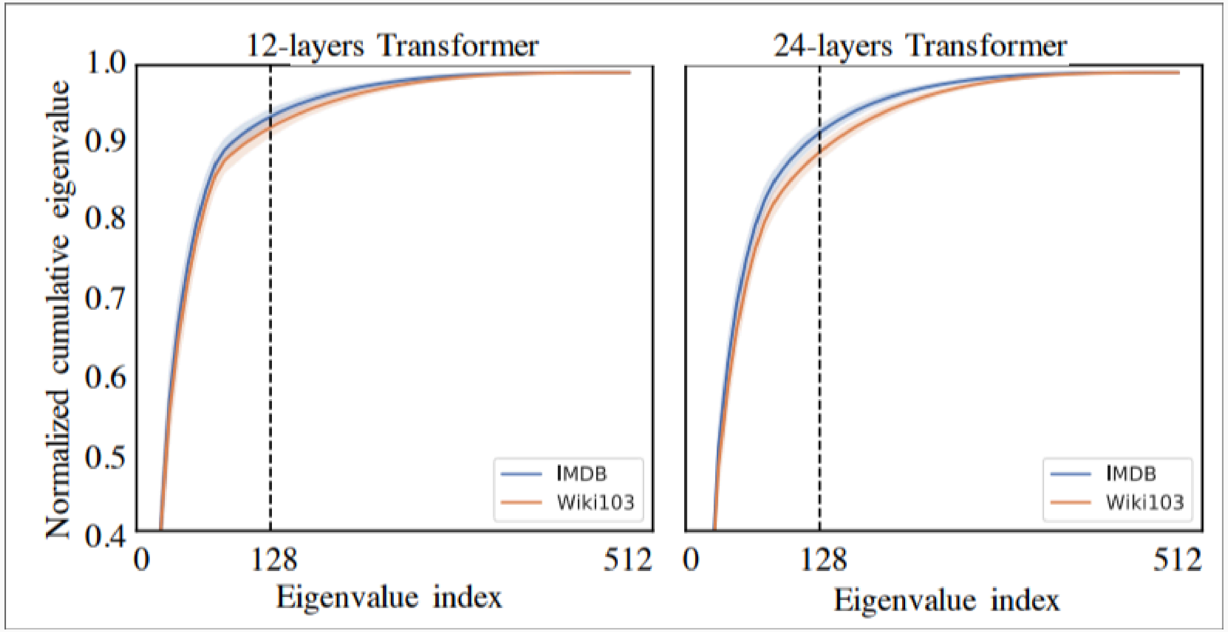

Figure 4: Eigenvalue plots (Wang, et al., 2020)

4. Questions?